

Topological Void Validation: Density-Controlled and Role-Aware Evaluation of Predicted Research Voids

Kris Pan
Intel Corporation
kris.pan@intel.com

Abstract. Topological Void Analysis (TVA) identifies candidate research voids in scientific knowledge spaces as geometric gaps between anchor-conditioned paper clusters. A natural validation question is whether future papers fill these predicted voids. We show that raw fill rate—whether any future paper appears near a void’s midpoint—fails as a validation target for three compounding reasons. First, it is confounded by local corpus density: a hot-zone baseline that samples high-density anchor-adjacent pairs outperforms TVA on raw fill (36% vs. 28%), but this hot-zone advantage disappears under a density-matched baseline: TVA shows only small paired differences (+1.3 to +3.0 pp across splits), with hierarchical cluster-bootstrap CIs all including zero and paired permutation $p > 0.35$; under calibrated thresholds, TVA and B2 collapse to near-parity (−0.7 to +0.7 pp). Second, it is limited by anchor observability: only 5–7% of future papers are anchor-eligible per split, so many unfilled voids should be treated as right-censored rather than invalid. Third, even the geometric threshold itself is density-confounded: a fixed threshold of 0.82 corresponds to a 35.7% average null occupancy rate, ranging from 1% to 99% across anchor-density cells. Under a null-calibrated threshold ($\tau_{\text{fill}} \approx 0.841$), fill rates collapse from $\sim 27\%$ to $\sim 0.7\%$, confirming that most raw fills are null-zone noise. Finally, role-aware epistemic classification shows that even anchor-eligible geometric fills rarely represent genuine void resolution: 59–76% are false positives across all groups and splits. We do not find statistically significant evidence that TVA produces more epistemic fills than density-matched baselines (TVA/B2 $\Delta = +3.0$ pp, hierarchical CI $[-2.0, +8.3]$ pp, $p = 0.36$). Instead, we show that the raw-fill validation protocol itself is unstable. These findings motivate a three-layer diagnostic framework and suggest that void validation requires calibrated, role-aware evaluation rather than raw geometric occupancy.

1 Introduction

Classical information retrieval asks: given a query, which existing documents are relevant? Topological Void Analysis (TVA) asks: given an anchor concept, what relevant knowledge is *absent*? TVA identifies candidate research voids as geometric gaps in an embedding-induced knowledge space—regions between anchor-conditioned paper clusters where expected connections, evidence, or

conceptual bridges are missing.

Validating predicted voids is not the same as validating a predicted link. A predicted link has a direct future target: the link appears or it does not. A predicted void has a richer failure mode: future literature may occupy the region geometrically while failing to bridge the source concepts epistemically. Unlike document retrieval or link prediction, predicted research voids have no established gold-standard benchmark: exhaustive

human annotation would require enumerating all candidate gaps, searching all future literature, and adjudicating whether each future paper plays a bridging role relative to specific source papers—a task infeasible at corpus scale. TVV addresses this by turning void validation into a staged, auditable protocol. The natural starting point—temporal fill rate—appears straightforward but is insufficient, and this paper studies why.

Contributions:

1. Raw temporal fill rate is strongly confounded by local historical density and research momentum.
2. A density-matched baseline shows that the apparent TVA disadvantage is not statistically significant (TVA/B2 $\Delta=+3.0$ pp, hierarchical CI $[-2.0, +8.3]$ pp, $p=0.36$).
3. Fixed cosine thresholds are not portable: 0.82 has 35.7% average null FPR; calibrated thresholds collapse raw fill to $\sim 0.7\%$.
4. Anchor eligibility exposes an observability limit: only $\sim 6\%$ of future papers are anchor-eligible; unfilled voids are right-censored, not invalid.
5. LLM-assisted role-aware audit suggests that geometric fill substantially overestimates epistemic resolution: 59–76% FP across all groups and splits.

Scope and framing. This paper is a *validation stress test*: we do not find statistically significant evidence that TVA outperforms density-matched baselines in epistemic fill prediction. The central contribution is methodological—showing that raw fill rate is an unstable validation proxy—rather than confirming TVA’s predictive superiority. Dynamic event-driven extensions are outside scope.

2 Background

2.1 Topological Void Analysis

TVA [11] operates on a corpus embedded with BGE-M3 [5] ($d = 1024$, cosine similarity). For

anchor q and source papers A, B , a candidate void satisfies four conditions: (C1) domain cohesion; (C2) calibrated marginality; (C3) sparse lexical bridge; (C4) midpoint vacancy. The midpoint is:

$$m(A, B) = \frac{\mathbf{a} + \mathbf{b}}{\|\mathbf{a} + \mathbf{b}\|} \in \mathbb{S}^{1023}.$$

2.2 Temporal Validation in LBD

Literature-Based Discovery validates predictions by checking whether future publications confirm predicted connections. Standard evaluation uses time-slicing and IR metrics (Recall@ K , Average Precision). TVV extends this by requiring that a future paper play a genuine epistemic bridging role, not merely appear geometrically nearby.

2.3 Research Momentum and Dense-Region Confounding

High-momentum regions near anchor topics attract disproportionate future publications regardless of void quality. Sampling baselines from these hot zones inflates apparent future fill rates, creating a density-confounded comparison.

3 Experimental Setup

3.1 Corpus and Temporal Splits

arXiv metadata (9 CS categories: OS, AR, PL, SE, DC, NI, CR, AI, LG), three rolling splits:

Table 1: Temporal splits.

Split	T	Val window	Train	Val
t3	2014	2015–2019	$\sim 29\text{k}$	$\sim 101\text{k}$
t4	2016	2017–2021	$\sim 35\text{k}$	$\sim 171\text{k}$
t5	2018	2019–2023	$\sim 51\text{k}$	$\sim 193\text{k}$

3.2 TVA Void Generation

10 anchor concepts $\times$ top-30 MMR voids = 300 candidate voids per split.

3.3 Baselines

B1 (hot-zone): 20 random pairs from top-300 anchor C1 pool per anchor. Biased toward high-density, high-momentum regions.

B2 (density-matched): For each TVA void, select the B1-style pair with midpoint local density $\rho(m) = \frac{1}{k} \sum \text{sim}(m, \text{kNN}_i(m))$ ($k = 20$) closest to the TVA void’s midpoint density.

3.4 TVV Three-Layer Protocol

We define three ordered indicator functions, applied as a sequential gate: if $E = 0$ then G and R are not evaluated; if $G = 0$ then R is not evaluated.

Local density. $\rho(V) := \frac{1}{k} \sum_{P \in \text{kNN}_k(m_V, \text{Train})} \text{sim}(P, m_V), \quad k=20.$

Anchor-eligibility threshold. $\tau_q = \max(Q_{80}^{\text{val}}, Q_{90}^{\text{train}})$ (computed per anchor and split).

Calibrated geometric threshold. Let $Z(m_0) = \max_{P \in \text{Val}_q} \text{sim}(P, m_0)$ for a null midpoint m_0 :

$$\tau_{\text{fill}}(q, \rho, t) = Q_{1-\alpha}(\{Z(m_0) : m_0 \in \text{Null}(q, \rho, t)\}),$$

where $\text{Null}(q, \rho, t)$ are density-matched null midpoints, calibrated set only (finite-sample conformal rank $k = \lceil (n_{\text{cal}} + 1)(1 - \alpha) \rceil$).

Three-layer indicators. Let $\text{Val}_q = \{P \in \text{Val}_t : \text{sim}(P, q) \geq \tau_q\}$ and $P_V^* = \arg \max_{P \in \text{Val}_q} \text{sim}(P, m_V)$.

$$E(V) := \mathbf{1}[|\text{Val}_q| \geq n_{\min}] \cdot \mathbf{1}[P_V^* \in \text{Val}_q] \quad (\text{anchor eligibility})$$

$$G(V) := \mathbf{1}[\text{sim}(P_V^*, m_V) \geq \tau_{\text{fill}}(q, \rho(V), t)] \quad (\text{geometric closure})$$

$$R(V) := \mathbf{1}[\text{role}(P_V^*, A, B, q) \in \{\text{TRUE-FILL, PARTIAL-FILL}\}] \quad (\text{epistemic role})$$

$$\text{Fill}(V) := E(V) \cdot G(V) \cdot R(V) \quad (1)$$

Proposition (monotone conservatism). For any void V : $\text{Fill}(V) \leq \text{Fill}_{\text{raw}}(V)$. Equality requires $E=G=R=1$ and $\tau_{\text{raw}} \leq \tau_{\text{fill}}(q, \rho, t)$. Each additional layer can only reduce apparent fill; no layer inflates it.

Algorithm 1 — TVV adjudication.

- 1: Compute Val_q ; if $|\text{Val}_q| < n_{\min}$: return $(0; E=0)$.
- 2: $P_V^* \leftarrow \arg \max_{P \in \text{Val}_q} \text{sim}(P, m_V)$.
- 3: Compute $\tau_{\text{fill}}(q, \rho(V), t)$ from calibration nulls.
- 4: If $\text{sim}(P_V^*, m_V) < \tau_{\text{fill}}$: return $(0; E=1, G=0)$.
- 5: $r \leftarrow \text{role-classifier}(P_V^*, A, B, q)$.
- 6: Return $(\mathbf{1}[r \in \{\text{TRUE, PARTIAL}\}]; E=1, G=1, R=r)$.

Why R requires text-based classification.

We investigated whether graph-theoretic surrogates can replace LLM-based role assessment for R . Specifically, we tested the change in effective resistance $\Delta R/R_{\text{before}}$ between A- and B-side communities before and after inserting candidate filler P into the k-NN graph (100 cases, t5). Two failure modes were observed: (1) in BGE-M3’s anisotropic 1024-dimensional space, local k-NN graphs are nearly fully connected, leaving $\Delta R \approx 0$ at machine precision; (2) on synthetic two-cluster graphs, inserting a bridge node can *increase* effective resistance due to altered flow-path distributions, making ΔR non-monotone with respect to bridging. These failures confirm that epistemic bridging cannot be reliably detected by local embedding topology, even when augmented with graph-spectral structure. The R layer’s reliance on text-based semantic assessment is therefore structurally necessary, not a methodological convenience.

4 Results

4.1 Finding 1: Raw Fill Rate Is Density-Confounded

After density matching, the apparent B1 hot-zone advantage disappears, but TVA does not signifi-

Table 2: TVA–B2 paired fill-rate differences (pp) under legacy $\tau=0.82$ and calibrated τ_{fill} . Hierarchical cluster bootstrap CI (anchor level, 1000 samples); p from paired anchor sign-flip permutation. All CIs include zero.

Split	Legacy $\tau=0.82$				Calibrated τ_{fill}		
	Δ_{pp}	95% CI	p		Δ_{pp}	95% CI	p
t3	+1.3	[−5.0, +9.0]	0.787		+0.7	[−1.0, +2.4]	0.412
t4	+2.4	[−1.4, +7.1]	0.352		−0.3	[−1.4, +0.8]	0.261
t5	+3.0	[−2.0, +8.3]	0.360		−0.7	[−2.0, +0.6]	0.141

cantly outperform B2. Under the legacy threshold, TVA–B2 differences are small and positive across splits (+1.3 to +3.0 pp), but hierarchical CIs all include zero and paired permutation tests are not significant ($p > 0.35$). Under calibrated thresholds, differences shrink to near-zero (−0.7 to +0.7 pp). Density-controlled temporal fill provides no statistically significant evidence of TVA superiority over B2.

This null result is expected under a strong counterfactual: B2 fixes the anchor and matches local midpoint density, estimating the background rate at which future papers occupy comparable regions. Parity shows that raw future occupancy is too low-base-rate and density-dependent to measure TVA’s added value as an ex-ante void-hypothesis generator—not that TVA lacks such value.

4.1.1 Fixed Thresholds Are Not Portable

Table 3: Null FPR@0.82 and τ_{fill} by density (t5).

Bucket	p50	τ_{fill}	FPR
Low	0.79–0.81	0.79–0.86	1–18%
Mid	0.80–0.83	0.82–0.86	5–80%
High	0.81–0.85	0.84–0.88	28–99%
Mean	—	0.841	35.7%

Under calibrated τ_{fill} , fill rates collapse: TVA: 27.3%→**0.7%**; B2: 24.3%→**0.7%**; lift: 1.00×

The 0.7% < 5% result is consistent with correct calibration: TVA C1–C4 conditions select sparser bridge regions than random density-matched pairs, so fewer reach the calibrated threshold.

Held-out calibration validity. To check that the calibration is not circular, we split null midpoints 80/20 (calibration/held-out). The held-out null FPR at τ_{fill} averages 6.3% across all anchor-density cells (target $\alpha=5\%$, $\Delta=+0.013$), indicating the calibration is slightly conservative but valid (Figure 1).




Figure 1: Null calibration curve: expected vs. observed held-out FPR across anchor-density cells (α range 0.01–0.20, t5). Points near the diagonal confirm that τ_{fill} is approximately calibrated. Mean held-out FPR at $\alpha=0.05$ is 6.3% (target 5%, $\Delta=+0.013$).

Most raw fills at $\tau = 0.82$ are null-zone occupancy, not significant geometric closure.




Figure 2: Calibrated threshold by density bucket (left) and null FPR at 0.82 (right). The fixed threshold 0.82 has 35.7% average null FPR, reaching 99% in high-density anchor cells. Calibrated $\tau_{\text{fill}} \approx 0.841$ is stable across all splits.

4.2 Finding 2: Anchor Eligibility Reveals Observability Limits

Table 4: Anchor exposure sample (t5).

Anchor	$\bar{s}_{im}$	τ_q	Pass	Fill
sched_opt	0.520	0.582	5.8%	40%
ebpf_obs	0.495	0.547	6.5%	0%
virt_hyp	0.498	0.555	7.7%	57%
Mean	0.503	0.563	6.2%	—

Only 5–8% of val papers are anchor-eligible per split. Voids with pass rate < 10% are flagged as *underexposed*: their unfill status reflects observability limits, not void invalidity. ebpf_obs (6.5% pass, 0% fill) illustrates that non-zero exposure does not guarantee eligible papers overlap predicted void neighborhoods.

Exposure-conditioned fill (censoring diagnostic). To assess whether unfilled voids are partly right-censored by observability, we bucket anchors by the number of anchor-eligible validation papers ($|Val_q|$) and report observed fill rates.

Table 5: Exposure-conditioned fill (t5). Fill increases monotonically with $|Val_q|$, indicating that non-fill partly reflects observability limits rather than false predictions. We therefore treat unfilled voids as right-censored, not as negatives.

Exposure bucket	Anchors	Mean $ Val_q $	Fill rate
Low	3	10,178	23.3%
Mid	4	11,794	26.7%
High	3	14,206	32.2%

Fill rates increase monotonically from 23.3% to 32.2% as anchor exposure increases. This does not prove that every unfilled void would be filled with more data, but it confirms that observed fill is partly a function of how much future literature is observable under each anchor. We do not interpret non-fill as a negative label.




Figure 3: Anchor eligibility pass rate (%) by anchor and split (hybrid gate). Rates are consistently low (5–8%), confirming corpus observability as a structural limit on temporal validation.

4.3 Role Classification Pipeline and Reliability

The LLM role classifier is the primary annotator; human review is an external plausibility check, not a gold standard (epistemic role is inherently interpretive).

Model and prompt. We use **claude-sonnet-4-5** (temperature 0.1) with a fixed nine-label rubric and structured JSON output. The system prompt instructs the model to judge the candidate paper solely relative to the anchor and source papers, ignoring general merit. Source group labels (TVA/B1/B2/near-miss) are *not shown* to the classifier (blind protocol).

Reliability checks. (i) *Self-consistency*: we re-run classification on a 100-case stratified subset three times; binary FP/non-FP agreement across runs exceeds 90%. (ii) *Human audit*: 60 cases stratified by source group and LLM label are reviewed by a blinded human auditor using the collapsed three-class rubric (void-resolution / boundary / FP). The auditor saw only anchor

text, A/B titles and abstracts, and candidate filler title and abstract. Binary FP/non-FP agreement between LLM and human auditor: $> 80\%$. Three-way collapsed label agreement: $\sim 65\%$, consistent with label boundary ambiguity in the PARTIAL-FILL / BOUNDARY-EXPANSION region.

Quantitative scope. All quantitative claims use collapsed three-way labels. Fine-grained nine-class labels are used only for qualitative case-study interpretation.

4.4 Finding 3: Geometric Fill Rarely Equals Epistemic Fill

Role labels are collapsed into three mutually exclusive classes for quantitative reporting. Table 6 defines each class operationally.

Table 6: Collapsed three-class role taxonomy. Fine-grained 9-class labels (TRUE-FILL, PARTIAL-FILL, IE, SE, BM, SN, RC, FP, UNCLEAR) are used for qualitative analysis only; all quantitative claims use the 3-class collapse.

Class	Operational criterion
Void-resolution	Paper explicitly bridges A and B under anchor q : proposes new method, mechanism, or synthesis connecting both source directions. Fine-grained: TRUE-FILL or PARTIAL-FILL.
Boundary activity	Paper extends one source direction or provides support, measurement, or naming relevant to the gap without bridging both sides. Fine-grained: IE, SE, BM, or SN.
False positive	Paper is geometrically close but does not address the anchor-conditioned void; topic, references, and framing are off-target. Fine-grained: FP.

An exploratory role audit shows 59–76% false positives across all groups and splits (pilot-scale; $n=204$ TVA gated cases across $t3+t4+t5$, role

Table 7: Role audit (1076 cases, $t3+t4+t5$; exploratory). VR=void resolution; FP=false positive. 95% Wilson CI in brackets (%).

Source	n	FP	FP CI	VR CI
TVA	204	71%	[65,77]	[1,6]
B1	158	75%	[67,81]	[0,2]
B2	202	67%	[60,73]	[1,6]
NM [†]	512	84%	[81,87]	[1,4]

[†]Near-miss ($\text{sim} \in [0.75, 0.82]$). CI in %.

labels from a single LLM). Near-miss controls show the highest FP rate (84%), confirming that geometric proximity carries some signal but is insufficient for void validation. TVA and B1 are indistinguishable; B2 is marginally better. These proportions are consistent but should be treated as directional, not confirmatory.




Figure 4: Epistemic role decomposition by source group ($t5$, 200 sampled cases). Most raw fills are geometric false positives regardless of method. Near-miss controls show the highest FP rate, confirming that geometric proximity is a weak but non-zero signal.

4.5 Landmark Three-Layer Diagnostic

As a positive-control diagnostic, we select five landmark bridge-type papers (pre-specified, widely cited, 2019–2023) and apply the full three-layer TVV predicate: geometric closure (G), anchor eligibility (E), and LLM-assisted bridge interpretability audit (R). Each landmark uses its publication-year cutoff as the train corpus bound-

ary.

Table 8: Landmark three-layer diagnostic. All five papers clear G and E; none clear R.

Landmark	G (>null p_{95})	E (anchor-elig.)	R (bridge audit)
vLLM / PagedAttention (2023)	Yes	Yes	No
MLIR (2020)	Yes	Yes	No
Devign (2019)	Yes	Yes	No
Ansor (2020)	Yes	Yes	Weak
FlashAttention (2022)	Yes	Yes	No

All five landmarks exceed the same-anchor null midpoint p_{95} threshold (G pass) and are anchor-eligible for at least one TVA anchor (E pass). However, LLM bridge audit (R) finds that the corresponding best-void source pairs are semantically incoherent: for example, vLLM’s best void pairs a GPU energy survey with a VANET routing paper, not OS paging with LLM serving.

This 5/5 G+E pass, 0/5 R pass pattern is precisely the failure mode TVV is designed to surface. It demonstrates that the three layers are not redundant: geometric proximity and anchor alignment are necessary screening signals but are insufficient without role-aware interpretation. The detector beeps at the right domain, but recovering the correct bridge pair requires the full adjudication pipeline.

4.6 Case Studies

The following three cases are selected from the t4/t5 role audit to illustrate why geometric proximity alone does not determine epistemic role. For each case we report the anchor direction, the two source-side communities, why the future paper is geometrically close, and why the role differs.

Case 1 – Partial Fill [9] (2019). Anchor: virtualization and cloud resource management. Source A lies on the VM resource allocation side; Source B lies on vehicular task-scheduling. The future paper studies multi-task offloading over vehicular clouds, combining resource allocation, task scheduling, and cloud/edge execution. Its embedding falls near the midpoint

($\text{sim}(P, m)=0.835$) because it shares vocabulary from both sides. However, its reference list contains 10+ task-offloading citations (B-side) but only one marginal VM reference (A-side). It approaches the bridge but does not fully connect both source communities. Role: PARTIAL-FILL.

Case 2 – Boundary Expansion [10] (2021).

Anchor: kernel scheduling and CPU affinity for heterogeneous systems. Source A represents the kernel-scheduler/CPU-affinity side; Source B represents heterogeneous-system optimization. The future paper proposes AI planning heuristics for heterogeneous HPC configuration and is geometrically close to the midpoint ($\text{sim}(P, m)=0.855$) through shared language of scheduling, planning, and resource optimization. Its entire bibliography (44 references) remains within the HPC optimization literature; there are no kernel scheduler, CPU-affinity, or OS-level references. It extends Source B’s direction without crossing toward Source A. Role: INCREMENTAL-EXTENSION.

Case 3 – False Positive [2] (2017).

Anchor: storage I/O path optimization and NVMe driver design. Despite having the *highest* midpoint similarity of the three cases ($\text{sim}(P, m)=0.875$), the paper studies energy-aware virtual network embedding. The high similarity is driven by generic vocabulary—“virtual”, “resource”, “energy”, “allocation”—shared with cloud infrastructure papers on both source sides, not by substantive connection to storage or NVMe research. Its 19 references contain no storage I/O, NVMe, block-layer, or anchor-domain citations. It is geometrically close but epistemically off-target. Role: FALSE-POSITIVE.

5 Discussion

5.1 Raw Fill Fails for Three Reasons

Raw fill conflates void quality with sociotechnical realization rate: (1) hot-zone baselines are biased by local momentum; (2) fixed thresholds are not portable across anchor-density regimes; (3) anchor-eligible future papers are sparse ($\sim 6\%$), right-censoring many unfilled voids.




Figure 5: Case-study role patterns. Case (c) has the *highest* $\text{sim}(P, m)$ yet is a false positive; Case (a) has lower similarity but partially bridges both source sides. Geometric proximity alone does not determine epistemic role.




Figure 6: Local density profile along the A-B bridge for the highest-MMR void in anchor sched_opt (t5). The density valley at $\lambda \approx 0.5$ marks the void midpoint region—two dense source communities with a less-occupied bridge.

5.2 B2 Is a Control, Not a Competitor

B2 is a strong counterfactual rather than a weak random baseline: it fixes the anchor and matches local midpoint density, estimating the background rate at which future literature occupies comparable regions without TVA’s structured void-selection criteria (C1–C4).

Under calibrated thresholds, TVA reaches parity with B2 but does not outperform it. We interpret this as a conservative result in a low-base-rate, corpus-limited regime: fill events are rare, both methods are exposed to the same small set of future papers, and the current corpus and validation windows may lack the statistical power to distinguish structured void discovery from density-matched background occupancy.

Parity with B2 is in fact the expected outcome under a properly controlled design. B2 matches the anchor and local density of TVA voids, so it estimates the background arrival rate of future papers in comparable regions. A large TVA advantage under this control would imply either strong method-specific predictive signal or residual confounding; we observe neither. Instead, both methods converge to the same low calibrated fill rate, indicating that significant geometric closure is primarily base-rate limited in this corpus.

B2 is intentionally conservative: it makes raw-fill superiority difficult to claim without evidence of signal beyond anchor and density. TVA does not provide that signal at this scale—but it was not designed to win on fill rate. Its value lies in turning a latent region into a named, inspectable, and actionable ex-ante hypothesis (A, B, q, m) before any future paper exists. Raw temporal fill rate is not a sensitive metric for that kind of value.

5.3 TVA as a Rapid Triage Instrument

The absence of a significant TVA–B2 fill-rate advantage should not be interpreted as an absence of practical utility. B2 is a strong *retrospective control*, not a discovery workflow: it estimates how often density-conditioned regions receive future occupancy. TVA, by contrast, produces a ranked list of inspectable void hypotheses (q, A, B, m) within minutes. Parity with B2 under dependence-aware tests is therefore nontrivial: TVA reaches the background future-occupancy rate of a strong matched control while also providing human-readable source-side structure for follow-up.

TVA is a *detector*, not an oracle. A high TVA signal is analogous to an anomaly reading: it prioritizes candidate locations for inspection, but does not certify epistemic closure. TVV provides the downstream adjudication protocol—calibrated geometric closure, anchor observability, and role-aware assessment—to determine whether a candidate void corresponds to a genuine epistemic gap.

Human-machine division of labor. The combinatorial search for candidate voids is machine-scale: a modest anchor neighborhood of 300 source papers induces 45 000 candidate pairs, and TVA generates 300 ranked void hypotheses across 10 anchors in minutes. Human experts are not designed for exhaustive enumeration at this scale, but are well-suited to adjudicating a compact, ranked list of interpretable (q, A, B, m) candidates. TVA therefore functions as a high-throughput triage step that compresses a machine-scale search space to a human-auditable hypothesis set.

Table 9: Division of labor: TVA performs high-throughput candidate generation; humans / LLM perform epistemic adjudication.

Task	Human expert	TVA/LLM
Enumerate source pairs	Infeasible at scale	Automatic
Rank candidate voids	Infeasible at scale	Automatic
Assess anchor alignment	Difficult	Automatic
Judge bridge meaning	Strong	Assisted, auditable
Validate technical merit	Strong	Exploratory

TVA should therefore be evaluated as a *rare-event triage instrument*. Because genuine epistemic closure is a low-base-rate event, most candidate voids will not become validated closures within any finite corpus window. The relevant measure is not absolute conversion rate but *enrichment*: does TVA concentrate observable future-bridging events into a smaller, inspectable candidate set relative to unguided or density-only search? Three rates matter at distinct levels: (i) *detection rate*—does TVA surface the right anchor domain? (ii) *interpretability rate*—does the A/B/q hypothesis make structural sense? and (iii) *realization rate*—does a future paper actually close the predicted gap? Only the first two are directly measurable from pre-publication corpora; realization is what TVV attempts to evaluate, and it is expectedly rare.

Table 10: Operational distinction between TVA and B2.

Property	TVA	B2
Purpose	Candidate void triage	Density-matched control
Output	(q, A, B, m) hypothesis	Matched midpoint
Interpretability	Source-side bridge	Not required
Runtime use	Ex-ante discovery	Retrospective evaluation
Validation role	Generates candidates	Estimates background rate

5.4 From Geometric Occupancy to Role Decomposition

Future papers near void midpoints decompose into true fills, boundary expansions, and false positives. The scarcity of true fills reveals that genuine epistemic closure is a much stricter condition than geometric occupancy. *The three layers are operational screens for disentangling three sources of evidence that raw fill rate conflates: geometric occupancy, observability, and epistemic role. They are not claimed to be necessary and sufficient conditions for scientific discovery.*

5.5 Descriptive, Not Evaluative

TVV classifies the type of validation evidence produced by future literature, not the intrinsic merit of papers. Boundary expansion and incremental extension are valuable scientific activities.

5.6 Dynamic Extension (Future Work)

These results suggest a natural extension: instead of asking whether a predicted void is eventually filled within a fixed window, one can classify each newly arriving paper by its pre-insertion topological impact on the existing knowledge space. This dynamic formulation is outside the scope of the present work.

5.7 Threats to Validity

(1) *Embedding dependence*: all metrics use BGE-M3. (2) *Threshold sensitivity*: calibration uses 200 null midpoints per cell. (3) *Corpus coverage*: fills outside arXiv are unobservable; unfilled voids are right-censored. (4) *LLM classification reliability*: role labels are produced by a single

LLM using a fixed rubric; we do not treat these as ground truth, since epistemic role is an inherently interpretive construct. To assess reliability we perform: (i) blind source-label removal (source group not shown to LLM), (ii) repeated classification on a 30-case subset for self-consistency, and (iii) a stratified human audit of coarse role categories (void-resolution vs. boundary vs. false positive) on 60 cases as an external plausibility check. The human audit is an auditor, not a gold standard: agreement between LLM and blinded human reviewer on the binary FP/non-FP split exceeds 80% on the audited subset. Fine-grained role agreement is lower ($\sim 65\%$), consistent with the inherent interpretability of the task. (5) *Domain specificity*: CS arXiv categories only.

6 Related Work

Validation of Predicted Voids. Prior retrospective evaluations in LBD [13, 8, 16] typically validate predicted connections by checking whether future publications instantiate an A–B association. Knowledge-graph completion similarly evaluates whether missing edges are recoverable under a fixed relation schema. These settings differ from void validation: a TVA void is not a predicted edge but an anchor-conditioned geometric vacancy whose future occupancy may represent true bridging, boundary expansion, or false-positive proximity. To our knowledge, prior work has not evaluated embedding-space voids under density-controlled, observability-aware, and role-aware temporal validation. TVV addresses this gap.

Literature-Based Discovery. Swanson’s ABC model [13] and subsequent LBD work [8, 16] validate predictions by checking future publication appearance. TVV extends this tradition by requiring epistemic bridging role, not merely geometric proximity.

Novelty and First-Story Detection. Topic Detection and Tracking [1] introduced first-story detection; novelty detection in IR [18] measures whether a document adds new information relative to a query set. Our first-filler intuition shares the event-ordering structure but operates

on knowledge structure rather than event reporting.

Structural Holes and Scientometrics. Burt’s structural hole theory [3] shows brokers spanning gaps accumulate informational advantage. Co-citation analysis [12] and our density-matched baseline directly control for hot-zone momentum.

Dense-Region Bias in Embeddings. Anisotropy in high-dimensional text embeddings [7] compresses cosine similarities into a narrow range. We show this anisotropy makes fixed thresholds non-portable in downstream void validation. Our use of “topological” is operational rather than homological [4, 6].

Scientific Claim Verification. Epistemic classification of scientific contributions—argumentative zoning [14], overclaiming detection [17]—provides precedents for our role-aware rubric. [15] introduced scientific fact verification as a retrieval task.

7 Conclusion

Temporal validation of predicted research voids requires separating geometric occupancy, density bias, anchor observability, and epistemic role. Under density matching, TVA’s apparent disadvantage disappears but is not statistically significant. Under null-calibrated thresholds, fill rates collapse to $\sim 0.7\%$ —most raw fills were null-zone noise. LLM-assisted role-aware classification suggests that genuine epistemic bridging is rare even among geometrically close, anchor-eligible future papers.

We propose a three-layer validation framework and leave two open questions for future work: (1) scalable role-aware validation without prohibitive LLM cost; (2) dynamic event-driven extensions in which each arriving paper updates knowledge-space state.

References

- [1] James Allan, Ron Papka, and Victor Lavrenko. On-line new event detection and

- tracking. In *Proceedings of SIGIR*, pages 37–45, 1998.
- [2] Amal S. Alzahrani and Ashraf A. Shahin. Energy-aware virtual network embedding approach for distributed cloud, 2017.
- [3] Ronald S. Burt. Structural holes and good ideas. *American Journal of Sociology*, 110(2):349–399, 2004.
- [4] Gunnar Carlsson. Topology and data. *Bulletin of the American Mathematical Society*, 46(2):255–308, 2009.
- [5] Jianlv Chen, Shitao Xiao, Peitian Zhang, Kun Luo, Defu Lian, and Zheng Liu. BGE M3-Embedding: Multi-lingual, multi-functionality, multi-granularity text embeddings through self-knowledge distillation, 2024.
- [6] Herbert Edelsbrunner and John Harer. *Computational Topology: An Introduction*. American Mathematical Society, 2010.
- [7] Kawin Ethayarajh. How contextual are contextualized word representations? Comparing the geometry of BERT, ELMo, and GPT-2 embeddings. In *Proceedings of EMNLP-IJCNLP*, pages 55–65, 2019.
- [8] Michael D. Gordon and Robert K. Lindsay. Toward discovery support systems: A replication, re-examination, and extension of Swanson’s work on literature-based discovery. *Journal of the American Society for Information Science*, 47(2):116–128, 1996.
- [9] Minghui Liwang, Zhibin Gao, Seyyedali Hosseinalipour, and Huaiyu Dai. Multi-task offloading over vehicular clouds under graph-based representation, 2019.
- [10] Suejb Memeti and Sabri Pllana. Optimization of heterogeneous systems with AI planning heuristics and machine learning: A performance and energy aware approach, 2021.
- [11] Kris Pan. Topological void analysis: A mathematical framework for systematic technical innovation discovery in knowledge spaces, 2026.
- [12] Henry Small. Co-citation in the scientific literature: A new measure of the relationship between two documents. *Journal of the American Society for Information Science*, 24(4):265–269, 1973.
- [13] Don R. Swanson. Fish oil, Raynaud’s syndrome, and undiscovered public knowledge. *Perspectives in Biology and Medicine*, 30(1):7–18, 1986.
- [14] Simone Teufel and Marc Moens. Summarizing scientific articles: Experiments with relevance and rhetorical status. *Computational Linguistics*, 28(4):409–445, 2002.
- [15] David Wadden, Shanchuan Lin, Kyle Lo, Lucy Lu Wang, Madeleine van Zuylen, Arman Cohan, and Hannaneh Hajishirzi. Fact or fiction: Verifying scientific claims. In *Proceedings of EMNLP*, pages 7534–7550, 2020.
- [16] Marc Weeber, Henny Klein, Lolkje T. W. de Jong-van den Berg, and Rein Vos. Using concepts in literature-based discovery: Simulating Swanson’s Raynaud–fish oil and migraine–magnesium discoveries. *Journal of the American Society for Information Science and Technology*, 52(7):548–557, 2001.
- [17] Dustin Wright and Isabelle Augenstein. Generating scientific claims for zero-shot scientific fact checking. In *Proceedings of ACL*, pages 2701–2720, 2022.
- [18] Yi Zhang, Jamie Callan, and Thomas Minka. Novelty and redundancy detection in adaptive filtering. In *Proceedings of SIGIR*, pages 81–88, 2002.

A Reproducibility Protocol

Corpus. arXiv metadata (title, abstract, categories, update_date) for categories cs.OS, cs.AR, cs.PL, cs.SE, cs.DC, cs.NI, cs.CR, cs.AI, cs.LG. Filtered by category; year assigned from update_date.

Temporal splits. Three non-overlapping splits. For split s with cutoff T_s : train = papers with year $< T_s$; val = papers with year $\in [T_s, T_s + 5)$.

Embedding. BGE-M3 (BAAI/bge-m3, 1024-dimensional, cosine similarity, offline/local mode). Papers embedded as “title abstract” concatenation.

TVA void generation. For each anchor q in Table 11 and each split: (1) embed anchor text with BGE-M3; (2) retrieve top-300 cosine-nearest train papers (C1 pool); (3) enumerate all pairs and apply C1–C4 conditions; (4) rank by MMR score and output top-30 voids. No anchor added or removed after validation runs began.

Table 11: TVA anchor list (pre-specified, fixed across all splits).

Key	Anchor phrase
sched_opt	kernel scheduler optimization and CPU affinity for multicore systems
mm_vm	virtual memory management and page reclamation in operating systems
ebpf_obs	eBPF programs for kernel tracing and system call observability
hws_w_x86	hardware-software co-design for x86 processor microarchitecture
net_io	high-performance network I/O and packet processing in the kernel
sync_lock	synchronization primitives and lock-free data structures
mem_alloc	memory allocator design and heap fragmentation in systems software
virt_hyp	virtualization and hypervisor design for cloud workloads
power_mgmt	dynamic power management and CPU frequency scaling
storage_io	storage I/O path optimization and NVMe device drivers

B1 baseline. For each anchor, sample 20 random pairs without replacement from the C1 pool (seed 42). No void conditions applied.

B2 density-matched baseline. For each TVA void V , compute midpoint density $\rho(m_V)$ (mean cosine similarity to $k=20$ nearest train neighbors).

Sample 300 B1-style pairs from the same anchor’s C1 pool (seed 99). Select the pair whose midpoint density is closest to $\rho(m_V)$ as the B2 counterpart.

Local density. $\rho(m) = \frac{1}{k} \sum_{i=1}^k \text{sim}(m, \text{kNN}_i(m, \text{train}))$, $k=20$, kNN over the train embedding matrix.

Anchor-eligibility threshold. $\tau_q = \max(Q_{80}(\text{val-anchor-sims}), Q_{90}(\text{train-anchor-sims}))$, computed per anchor and split.

Null-calibrated threshold. For each cell (anchor q , density bucket $b \in \{\text{low, mid, high}\}$, split t): (1) generate $n=1000$ null midpoints from C1 pool; (2) split 50/50 calibration/held-out; (3) compute $Z(m_0) = \max_{P \in \text{Val}_q} \text{sim}(P, m_0)$ for each null; (4) $\tau_{\text{fill}} = s_{(k)}$ where $k = \lceil (n_{\text{cal}} + 1)(1 - \alpha) \rceil$ (finite-sample conformal quantile), $\alpha=0.05$.

Statistical tests. Primary inference uses *hierarchical cluster bootstrap* (outer: resample anchors with replacement; inner: resample voids within each anchor) for TVA–B2 difference in percentage points. Secondary check: *anchor-level sign-flip permutation test* (randomly flip sign of per-anchor mean TVA–B2 difference; $n_{\text{perm}}=10000$). Both report two-sided p -values.

Role annotation. Model: Claude Sonnet 4.5, temperature 0.1. Input: anchor text, source A title+abstract (truncated), source B title+abstract, candidate filler title+abstract. Output: JSON with role label, confidence, and six ordinal scores. Scores group labels are not shown (blind classification). Final labels use the collapsed three-class taxonomy (Table 6).

Prompt and full rubric text available at: <https://github.com/krispan-intel/deepthought>.

Code and data availability. Experiment scripts, prompts, and configuration files are available at <https://github.com/krispan-intel/deepthought> (branch: dynamic.tva). Derived labels for the role-classified 1076 cases will be released upon acceptance.